Table of Contents

1 Measurements and results

1.1 Model comparisons

As mentioned in section **??** we used Keras Tuner model to find hyperparameters that would yield the highest accuracy. Instead of hard-coding hyperparameters when building a model with Keras API, we defined a search space of possible values with `HyperParameter` class and used that as a hyperparameter.

We then passed the created model to a `RandomSearch` class, with few other parameters such as batch size, number of epochs and maximum number of trials. We then started hyperparameter search, which means that Keras Tuner was randomly picking a set of hyperparameters and training a model with them. This process was repeated for a trial number of times. Accuracy and hyperparameters that were used while training every module were saved to a log file for later analysis.

After training a number of different models we handpicked a few of them and compared them. Comparison of equivalent model trained in Edge Impulse studio was also done.

1.1.1 Hyperparameter search space and results analysis

General structure of CNN model was already described in section **??** and in Figure **??**. We decided to search for following hyperparameters: number of filters in all three convolutional layers (can be different for each layer), filter size in all three convolutional layers (same for all layers), size of dense layer, dropout rate and learning rate. Possible values of hyperparameters (also known as hyperparameter search space) are specified in table 1.1.

Hyperparameter	Set of values
FilterNum1	From 16 to 80, with a step of 8
FilterNum2	From 16 to 80, with a step of 8
FilterNum3	From 16 to 80, with a step of 8
FilterSize	3 x 3 or 3 x 4
DenseSize	From 16 to 96, with a step of 8
DropoutRate	From 0.2 to 0.5, with a step of 0.05
LearningRate	0.0001 or 0.0003
Random search variable	**value**
EPOCHS	25
BATCH_SIZE	100
MAX_TRIALS	300

Table 1.1: First hyperparameter search space

Search space of `filter_numX`, `dense_size` and `dropout_rate` hyperparameters was chosen based on initial training tests and various models that were trained on similar data. Value of `filter_size` is usually 3 x 3, however all example ML projects were training on image date of same dimensions. We wanted to test how would a filter with same ratio of dimensions as image data (3 x 4 and 60 x 80 respectively) perform. Hyperparameter `learning_rate` was chosen heuristically, we saw that higher values, such as 0.001 or 0.003, would leave model's accuracy stuck at suboptimal optima, from where it could not be improve anymore.

We also had to set 3 variables that directly affected how long will random search last. From initial tests we saw that models usually reached maximum possible accuracy around 20^{th} epoch, to give some headroom we set the number of epochs to 25. We kept batch size relatively small, at 100, which meant that weights would get updated regularly. Hyperparameter `MAX_TRIALS` had the biggest impact on the training time, we set it to 300.

Training lasted for about 12 hours. After it was done we compiled a list of all 300 trained models and their different hyperparameter values, number of parameters and accuracies, part of it can be seen in Table 1.2.

It is important to keep in mind that we are dealing with imbalanced dataset, where

Hyperparameter	FilterNum1	FilterNum2	FilterNum3	DenseSize	DropoutRate	FilterSize	LearningRate	Number of parameters	Accuracy[%]
Model ID									
0a	72	80	64	72	0.4	3x4	0.0003	1.514,400	98.35
1a	32	40	72	56	0.35	3x4	0.0001	1.260,332	98.31
2a	40	48	32	64	0.35	3x4	0.0001	656,797	98.31
3a	56	16	48	72	0.4	3x4	0.0001	1,057,924	98.28
4a	80	64	40	96	0.45	3x4	0.0003	1,245,788	98.28
5a	64	24	72	88	0.45	3x3	0.0001	1,931,356	98.28
6a	64	56	40	80	0.35	3x3	0.0001	1,013,556	98.24
7a	16	40	64	88	0.35	3x3	0.0003	1,719,108	98.24
8a	48	64	32	64	0.35	3x4	0.0003	676,884	98.24
9a	72	48	56	80	0.3	3x4	0.0001	1,419,172	98.24
91a	72	48	56	40	0.35	3x3	0.0003	728,324	98.00
92a	48	24	64	64	0.35	3x4	0.0001	1,262,092	98.00
93a	24	48	32	72	0.4	3x3	0.0001	716,076	98.00
94a	32	48	72	32	0.25	3x3	0.0001	736,732	98.00
95a	64	24	64	48	0.3	3x4	0.0001	959,628	98.00
96a	16	32	72	80	0.25	3x4	0.0001	1,762,508	98.00
97a	72	56	40	56	0.45	3x4	0.0003	748,580	98.00
98a	32	24	24	48	0.35	3x3	0.0001	358,308	98.00
99a	48	16	40	40	0.45	3x3	0.0003	493,412	98.00
100a	24	72	64	40	0.45	3x3	0.0003	844,684	98.00
191a	64	56	16	52	0.4	3x3	0.0001	386,996	97.76
192a	48	40	24	24	0.4	3x4	0.0001	208,172	97.73
193a	56	64	72	24	0.25	3x4	0.0003	617,692	97.73
194a	48	72	48	32	0.25	3x4	0.0003	544,652	97.73
195a	72	56	24	56	0.25	3x4	0.0003	469,012	97.73
196a	72	48	72	40	0.3	3x3	0.0003	927,252	97.73
197a	80	16	32	80	0.25	3x3	0.0001	785,380	97.73
198a	56	24	16	88	0.25	3x3	0.0001	438,996	97.73
199a	56	24	16	88	0.25	3x3	0.0001	438,996	97.73
295a	48	32	64	16	0.5	3x4	0.0001	351,012	95.87
296a	40	24	56	24	0.5	3x4	0.0001	431,572	95.77
297a	56	16	80	16	0.2	3x4	0.0001	411,020	95.63
298a	24	16	48	24	0.5	3x4	0.0001	359,924	94.46
299a	40	48	56	16	0.35	3x3	0.0003	310,860	82.86

Table 1.2: Partial results of first random search of hyperparameters

82.86 % of our validation data are elephant images. Simply classifying all images as elephant class would yield 82.86 % accuracy, which sounds high, although it is actually not helpful.

After analyzing results we came to several conclusions:

1. We saw that almost all trained models, except of the last one, achieved accuracy above 90 %. This proved that the general architecture of the model was appropriate for the problem.

2. We could not see any visible correlation between a specific choice of a certain hyperparameter and accuracy. This shows that selection of hyperparameters is really a non-heuristic task.

3. Filter of size 3 x 4 did not perform significantly better compared to one with size 3 x 3.

4. There is a weak correlation between number of parameters (model's complexity) and accuracy. Although eight of top ten models have more then 1 million parameters, models with IDs 2 and 8 have almost half of the parameters, but still perform well.

5. First 200 models cover an accuracy range of 0.62 %. However inside of this range model number of parameters varies hugely, for example, model with ID 192a has more than 8 times less parameters than model than model with ID 96a, although the difference in accuracy (0.27 %) is negligible.

As we realized that we do not need complex models to achieve high accuracy on our training data, we decided to run the random search of hyperparameters again. We decided to lower the maximum and minimum numbers of filters and size of dense layer. We also decreased the step from 8 to 2. We decided to lower the bottom boundary of `DropoutRate` from 0.2 to 0.0, which means that some models will not be using dropout layer at all. We expected that training without dropout layer would produce suboptimal results, however we wanted to test it. Redefined search space for second random search can be seen in Table 1.6 We increased the number of `MAX_TRIALS` from 300 to 500, as we were expecting that more models will end

Hyperparameter	Set of values
FilterNum1	From 4 to 48, with a step of 2
FilterNum2	From 4 to 48, with a step of 2
FilterNum3	From 4 to 48, with a step of 2
FilterSize	3 x 3 or 3 x 4
DenseSize	From 4 to 48, with a step of 2
DropoutRate	From 0.0 to 0.5, with a step of 0.05
LearningRate	0.0001 or 0.0003
Random search variable	**value**
EPOCHS	25
BATCH_SIZE	100
MAX_TRIALS	500

Table 1.3: Second hyperparameter search space

up underfitting and also because there would be more possible options because of smaller step size.

Partial table of results of random hyperparameter search can be seen in Table 1.5.

Some observations:

1. We can see that the accuracy of the best model 0b compared to the best model 0a from previous search is only 0.21 % lower, although it has about 5 times less parameters.

2. Although that it might seem that FilterSize of 3 x 4 yields best results, we did not saw a strong bias towards on or another option after manually analyzing best 30 models.

3. We can see that worst five models have the same accuracy of 82.86 %, same as the worst performing model from first random search. There are 82.86 % images of elephants in validation class, which means that model probably assigned all validation images to elephant class and was satisfied with achieved accuracy.

4. We can see that the model with ID 296b has quite low number of parameters,

Hyperparameter	FilterNum1	FilterNum2	FilterNum3	DenseSize	DropoutRate	FilterSize	LearningRate	Number of parameters	Accuracy[%]
Model ID									
0b	40	20	20	48	0.25	3x4	0.0001	304,216	98.14
1b	44	10	28	42	0.2	3x4	0.0003	362,264	98.14
2b	18	38	26	38	0.1	3x4	0.0003	316,956	98.11
3b	46	34	28	40	0.35	3x4	0.0003	367,056	98.07
4b	26	36	36	34	0.15	3x4	0.0003	394,568	98.04
95b	20	16	34	40	0.3	3x3	0.0003	416,230	97.62
96b	46	42	28	32	0.4	3x3	0.0003	297,466	97.62
97b	30	26	30	34	0.2	3x3	0.0001	320,570	97.59
98b	20	10	16	46	0.45	3x3	0.0003	224,500	97.59
99b	32	20	32	26	0.25	3x4	0.0003	265,562	97.59
195b	28	16	40	24	0.1	3x3	0.0001	298,252	97.31
196b	44	30	32	20	0.3	3x4	0.0003	220,098	97.31
197b	46	40	10	40	0.1	3x3	0.0001	140,874	97.31
198b	12	28	40	32	0.4	3x4	0.0003	401,860	97.31
199b	36	38	6	40	0.4	3x3	0.0003	86,972	97.31
295b	20	8	34	26	0.3	3x3	0.0003	269,464	96.90
296b	18	16	10	20	0.3	3x4	0.0003	65,740	96.87
297b	8	22	28	16	0.1	3x3	0.0001	141,742	96.87
298b	18	28	14	32	0.25	3x4	0.0001	145,592	96.87
299b	28	40	20	12	0.25	3x3	0.0001	89,684	96.87
395b	10	20	12	30	0.0	3x3	0.0001	112,246	96.87
396b	24	24	46	18	0.2	3x3	0.0003	263,924	96.14
397b	6	18	12	24	0.4	3x4	0.0001	90,520	96.11
398b	44	22	6	32	0.45	3x4	0.0001	71,564	96.11
399b	12	16	10	28	0.25	3x4	0.0001	88,550	96.08
495b	32	14	40	8	0.5	3x4	0.0003	109,334	82.86
496b	36	38	22	6	0.2	3x3	0.0003	59,890	82.86
497b	42	30	22	6	0.4	3x3	0.0003	57,386	82.86
498b	4	4	20	12	0.4	3x3	0.0003	72,992	82.86
499b	32	36	36	4	0.15	3x3	0.0001	65,648	82.86

Table 1.4: Partial results of second random search of hyperparameters

Hyperparameter	FilterNum1	FilterNum2	FilterNum3	DenseSize	DropoutRate	FilterSize	LearningRate	Number of parameters	Accuracy[%]
Model ID									
172b	42	44	8	14	0.1	3x4	0.0001	60,672	97.38
230b	34	6	4	40	0.2	3x3	0.0003	50,606	97.18
265b	6	46	6	24	0.3	3x3	0.0003	48,404	97.01
270b	36	20	12	10	0.15	3x3	0.0003	45,086	97.01
304b	46	6	10	12	0.2	3x4	0.0003	40,710	96.87
338b	4	18	6	10	0.05	3x4	0.0003	20,290	96.63
454b	10	14	8	6	0.05	3x4	0.0001	17,610	94.18
460b	6	28	4	8	0.1	3x4	0.0003	13,114	93.60

Table 1.5: Partial results of second random search of hyperparameters

only 65,740.

After making a skim analysis of results, we decided to go through the whole list of trained results and mark down those with extremely low number of parameters, they are presented in table **??**.

1.1.2 Comparison of selected, re-trained models

Two random searches gave us a large amount of different models to choose from. In every other ML application where the execution time would not be a constraint, we could simply take the best performing model and be done with it. In our case we need to make a trade off between model accuracy and execution speed. In order to do that we need to understand what a difference of few percents of accuracy actually means.

As we are dealing with imbalanced dataset, where 82.86 % of our validation data are elephant images, confusion matrix, precision and recall metrics[1]will give us better idea how well is the model performing.

---

[1]Precision tells us what percentage of data points in a specific predicted class actually fall into that class. Recall tells us what percentage of data points of a class were actually predicted correctly [1].

Models	Model 0a	Model 460b
Metrics		
accuracy[%]	98.18	93.4
number of parameters		
Precision of elephant class	99.22	97.8
Precision of human class	96.92	80.31
Precision of cow class	90.99	69.37
Precision of nature/random class	77.42	40.32
Recall of elephant class	99.29	97.39
Recall of human class	93.2	85.57
Recall of cow class	94.39	81.91
Recall of nature/random class	87.27	28.09

Table 1.6: Second hyperparameter search space

For our initial comparison we chose best performing model from first search (ID 0a, accuracy 98.35 %) and the model with the least number of parameters from second search (ID 460b, accuracy, 93.60 %). Figure 1.1 shows two confusion matrices of those two models, Table **??** shows comparisons between precision and recall metrics.


Figure 1.1

1.1.3 Comparison with Edge Impulse model

Choose one, or few model configurations and make a model in EI.

Describe how did you build Edge Impulse model, maybe mention transfer learning

(might be extra work, you did not write about theory)

1.2 On device performance testing

Describe setup

To profile execution of our code we first wrote a timer driver based on a Arm's systick timer, however we later decided to use data watch trigger (DWT). DWT does not use interrupts, therefore it does not introduce overhead of calling interrupt routines like systick timer does.

how are you timing,

1.2.1 Comparison of different optimization options

1.2.2 Comparison of performance of selected models

Transfer learning could be interesting here, bigger number of parameters and takes less time.

1.3 Power profiling of an embedded early warning system

- [ ] Power consumption test of whole setup, PIR wakes up wisent, wisent turns on stm32f7 and flir, which makes a picture, does inference, reports result and wisent sends the result. Otii image of consumption with marked sections.(shouldnt be hard)

1.3.1 Battery life estimations

Based on numbers and different scenarios estimate how long would this last with different batteries.

Bibliography

[1] Geron, A. *Hands-on machine learning with Scikit-Learn and TensorFlow: concepts, tools, and techniques to build intelligent systems, 2nd edition.* O'Reilly Media, Sebastopol, CA, 2019.